

Project 1: MESI Cache Coherence

Sangeeth Thayaaparan
RIN: 662037988

February 26, 2026

Tools: Quartus Prime 20.1.1, ModelSim - Intel FPGA Starter Edition 2020.1

Abstract

This report goes over the design, implementation, and verification of a cache-coherent memory system, specifically implementing the MESI coherence protocol. The system consists of 4 processors, each with its own L1 cache, a shared L2 cache, and DRAM. The CPU testbench generates memory requests to test various scenarios (hits, misses, writebacks) and the cache controller FSM manages the state transitions to maintain coherence. More information regarding the design and verification can be found in the following sections.

Contents

1 Spec and Rationale	3
2 Microarchitecture and Cycle Behavior	3
2.1 Full System Architecture Diagram	3
2.2 Module Descriptions	4
3 Verification and Results	10
3.1 Test 1: Memory Read	11
3.2 Test 2: Memory Write	12
3.3 Test 3: Cache Coherence (Snooping + MESI Transitions)	13
3.4 Test 4: Arbitration (Round-Robin Fairness Under Contention)	14
3.5 Test 5: Memory Eviction and Writeback	15
A Test Vector Files	17
A.1 Test 1 Vectors (Memory Read)	17
A.2 Test 2 Vectors (Memory Write)	17
A.3 Test 3 Vectors (Cache Coherence)	17
A.4 Test 4 Vectors (Arbitration)	18
A.5 Test 5 Vectors (Eviction/Writeback)	18
B Reproducibility Notes	19

List of Figures

1	Full Cache System Block Diagram	3
2	Full System Architecture Diagram (created by author).	4
3	Round-robin arbitration illustration (adapted from [3]).	9
4	MESI state transition diagram (from [1], credited to [2]).	11
5	Test 1 waveform showing memory read operations, cache hits/misses, and MESI state transitions.	12
6	Test 2 waveform showing memory write operations, cache hits/misses, MESI state transitions, and bus transactions.	13
7	Test 3 waveform showing coherence snoops, shared-line formation, invalidation, and ownership transfer.	14
8	Test 4 waveform showing concurrent requests and round-robin bus grant sequencing.	15
9	Test 5 waveform showing conflict misses, dirty evictions, writeback transactions, and reload behavior.	16

List of Tables

1	Test 1 (Memory Read): stimulus and expected cache/coherence behavior	11
2	Test 2 (Memory Write): stimulus and expected cache/coherence behavior	12
3	Test 3 (Cache Coherence): stimulus and expected cache/coherence behavior	13
4	Test 4 (Arbitration): concurrent request groups and expected arbiter behavior	14
5	Test 5 (Eviction/Writeback): stimulus and expected behavior	15

1 Spec and Rationale

The L1 caches are almost an exact copy of the cache from project 0, the specs being:

- Cache Size: 512 blocks (8 KB total)
- Block Size: 128 bits (4 words)
- Address Breakdown:
 - Tag: 19 bits (bits [31:13])
 - Index: 9 bits (bits [12:4])
 - Offset: 2 bits (bits [3:2])

For simplicity reasons, the L2 cache is the exact same as the L1 cache. The DRAM is a simple memory model that can read/write 128-bit blocks (4 words) with single-cycle latency.

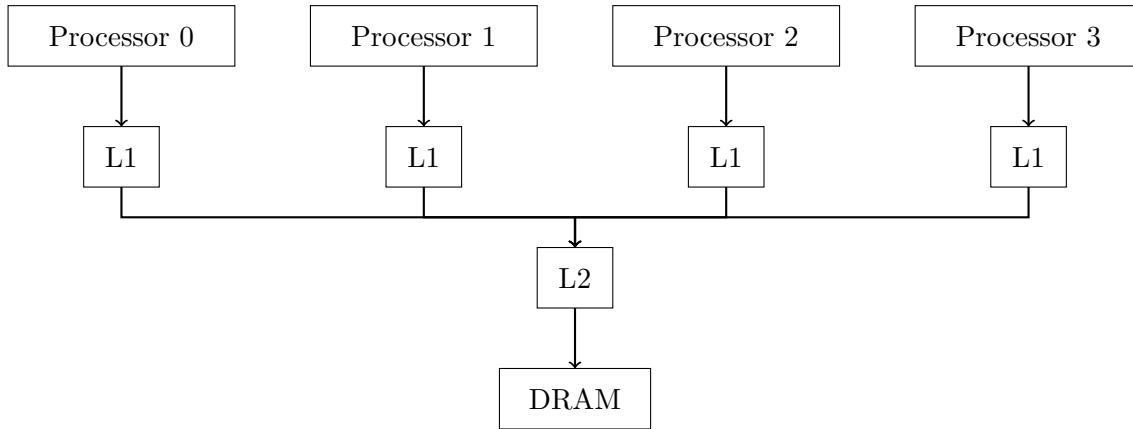


Figure 1: Full Cache System Block Diagram

2 Microarchitecture and Cycle Behavior

2.1 Full System Architecture Diagram

The full system architecture was designed using KiCAD. While each module doesn't have too much to it (the inside of each subsheet just contains inputs and outputs, and the L1 cache contains multiple variants of the MESI FSM), this schematic was mostly used to help design the verilog code. Specifically picking and choosing the names of each wires and how the modules are connected together. While the final verilog code modifies some of the names (as there are various repeat names for different lines and buses), the overall structure of the system was based off this schematic.

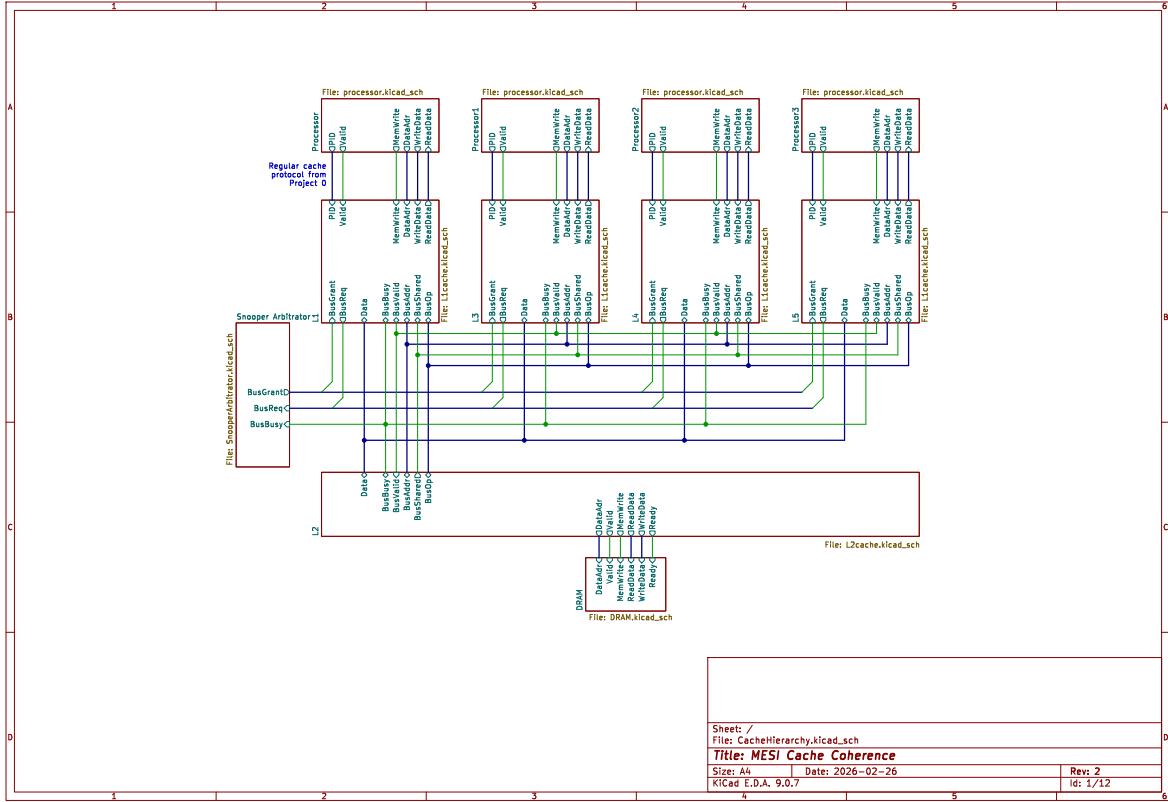


Figure 2: Full System Architecture Diagram (created by author).

2.2 Module Descriptions

2.2.1 Module: CPU

Behavior: The cpu module isn't actually implemented as it's own module in the verilog code. Rather, the testbench generates the CPU requests directly. The CPU generates memory requests with an address, data, read/write signal, and a *valid* signal. It waits for the cache to assert *Ready* before issuing the next request. Correctness means: all test vectors are issued in order and the CPU properly waits for completion.

Key signals:

- **Valid** (1 bit, output): Indicates a new request is being issued
- **DataAddr** (32 bits, output): Address for the memory operation
- **WriteData** (32 bits, output): Data to write (for write operations)
- **MemWrite** (1 bit, output): 1 for write, 0 for read
- **Ready** (1 bit, input): Asserted by cache when request is complete

Key State:

- **testvectors[0:99]**: Array of test vectors loaded from file, each containing {address, data, write}
- **vectornum**: Index of the current test vector being issued

Implementation:

```
1 typedef struct {
2     logic [31:0]    addr;
3     logic           write;
4     logic [1:0]     l1_id;
5     logic [31:0]    data;
6     int             cycle;
7 } test_vector_t;
```

Listing 1: Test Vector Format (from testbench)

```
1 for (int cycle = 0; cycle <= max_cycle; cycle++) begin
2     test_vector_t pending_vectors[$];
3     pending_vectors = vectors_by_cycle[cycle];
4
5     if (pending_vectors.size() > 0) begin
6         for (int l1 = 0; l1 < 4; l1++) dut.l1_valid[l1] = 1'b0;
7
8         foreach (pending_vectors[i]) begin
9             issue_request(pending_vectors[i]);
10            end
11
12        wait_for_all_ready(pending_vectors, max_cycle, cycle);
13    end
14 end
```

Listing 2: Cycle-Based Request Issue (from testbench)

```
1 task automatic wait_for_all_ready(ref test_vector_t pending_vectors[$],
2                                     input int max_cycle,
3                                     input int current_cycle);
4     int done_count = 0;
5     int total_count = pending_vectors.size();
6     bit ready_mask[];
7
8     ready_mask = new[total_count];
9     @(posedge clk);
10
11    while (done_count < total_count) begin
12        done_count = 0;
13        foreach (pending_vectors[i]) begin
14            if (!ready_mask[i] && dut.l1_ready[pending_vectors[i].l1_id]) begin
15                ready_mask[i] = 1'b1;
16                dut.l1_valid[pending_vectors[i].l1_id] = 1'b0;
17            end
18            if (ready_mask[i]) done_count++;
19        end
20        if (done_count < total_count) @(posedge clk);
21    end
22 endtask
```

Listing 3: Per-Request Completion Handling (from testbench)

2.2.2 Module: L1 cache

Behavior: The L1 module is a direct-mapped private cache for each processor with MESI coherence states (INVALID/SHARED/EXCLUSIVE/MODIFIED). It accepts CPU read/write requests, serves hits locally, and issues bus transactions on misses or upgrades (BusRd, BusRdX, BusUpgr). It also snoops peer transactions and performs coherence state transitions plus invalidations as needed.

Key signals:

- Valid, MemWrite, DataAddr, WriteData: CPU-side request interface
- ReadData, Ready, CacheHit: CPU-side response/status interface
- BusReq, BusGrant: arbitration handshake
- Data, BusAddr, BusOp, BusValid, BusShared, BusBusy: shared coherence bus

Key State:

- SRAM[511:0]: cache lines with MESI state, dirty bit, tag, and 128-bit block data
- ctrl_state: transaction controller FSM (IDLE, REQUEST_BUFFERED, MISS_PENDING, BUS_ACTIVE, WRITEBACK_PENDING)
- Buffered request registers (req_addr, req_wdata, req_write, req_valid) and writeback buffer (wb_addr, wb_data, wb_pending)

Implementation:

```

1  typedef enum logic [1:0] {
2      INVALID,
3      SHARED,
4      EXCLUSIVE,
5      MODIFIED
6  } mesi_t;
7
8  typedef enum logic [2:0] {
9      IDLE,
10     REQUEST_BUFFERED,
11     MISS_PENDING,
12     BUS_ACTIVE,
13     WRITEBACK_PENDING
14 } ctrl_t;
```

Listing 4: L1 MESI + Controller States

```

1 assign cache_hit_comb = req_valid &&
2                     (SRAM[req_index].mesi_state != INVALID) &&
3                     (SRAM[req_index].tag == req_tag);
4
5 REQUEST_BUFFERED: begin
6     if (cache_hit_comb) begin
7         CacheHit = 1'b1;
8         if (req_write && SRAM[req_index].mesi_state == SHARED) begin
9             ctrl_next = MISS_PENDING; // upgrade via bus
10            BusReq = 1'b1;
11        end else begin
12            Ready = 1'b1;
```

```

13     ctrl_next = IDLE;
14   end
15 end else begin
16   ctrl_next = (SRAM[req_index].dirty && SRAM[req_index].mesi_state != 
17     INVALID)
18   ? WRITEBACK_PENDING : MISS_PENDING;
19 end
end

```

Listing 5: L1 Hit Detection and Hit/Miss Path

2.2.3 Module: L2 cache

Behavior: The L2 module is a shared inclusive cache and coherence responder on the bus. It snoops BusRd/BusRdX/BusUpgr/Writeback, returns data on L2 hits, triggers DRAM reads on L2 misses, and installs/updates returned lines. It also tracks sharer information and can assert BusShared.

Key signals:

- Data, BusAddr, BusOp, BusValid, BusShared, BusBusy: shared coherence bus interface
- dram_valid, dram_write, dram_addr, dram_wdata: DRAM request interface
- dram_rdata, dram_ready: DRAM response interface

Key State:

- L2_SRAM[511:0]: valid/tag/sharer mask/data per cache line
- DRAM miss tracking: dram_pending, dram_pending_tag, dram_pending_index, dram_pending_busadr

Implementation:

```

1 typedef struct packed {
2   logic      valid;
3   logic [18:0] tag;
4   logic [3:0]  l1_sharers;
5   logic [127:0] data;
6 } l2_block_t;
7
8 if (BusOp === 3'b000 || BusOp === 3'b001) begin
9   if (L2_SRAM[bus_index].valid && L2_SRAM[bus_index].tag == bus_tag) begin
10     data_out = L2_SRAM[bus_index].data;
11     data_oe = 1'b1;
12     busvalid_out = 1'b1;
13   end else begin
14     dram_valid = 1'b1;
15     dram_write = 1'b0;
16     dram_addr = {BusAddr[31:4], 4'b0000};
17   end
18 end

```

Listing 6: L2 Line Format and DRAM Miss Path

```

1 if (dram_pending && dram_ready) begin
2   L2_SRAM[dram_pending_index].valid <= 1'b1;
3   L2_SRAM[dram_pending_index].tag <= dram_pending_tag;
4   L2_SRAM[dram_pending_index].data <= dram_rdata;
5   dram_pending <= 1'b0;
6 end
7
8 if (BusOp === 3'b011 && !$isunknown(Data)) begin // Writeback
9   L2_SRAM[bus_index].valid <= 1'b1;
10  L2_SRAM[bus_index].tag <= bus_tag;
11  L2_SRAM[bus_index].data <= Data;
12 end

```

Listing 7: L2 DRAM Fill and Bus Update

2.2.4 Module: dram

Behavior: The `dram` module models backing main memory using a 32-bit word array and supports 128-bit block transfers (4 consecutive words) for cache-line fills and writebacks. On `Valid`, it performs either a block write or block read and asserts `Ready`.

Key signals:

- `Valid`: transaction request
- `MemWrite`: write/read select
- `DataAddr`: starting word address
- `WriteDataBlock`, `ReadDataBlock`: 128-bit block data interface
- `Ready`: transaction complete

Key State:

- `DRAM[1048575:0]`: 1,048,576-word memory array (32-bit words)

Implementation:

```

1 always @(posedge clk) begin
2   if (Valid) begin
3     if (MemWrite) begin
4       DRAM[DataAddr]      <= WriteDataBlock[31:0];
5       DRAM[DataAddr + 1]  <= WriteDataBlock[63:32];
6       DRAM[DataAddr + 2]  <= WriteDataBlock[95:64];
7       DRAM[DataAddr + 3]  <= WriteDataBlock[127:96];
8     end else begin
9       ReadDataBlock[31:0]   <= DRAM[DataAddr];
10      ReadDataBlock[63:32] <= DRAM[DataAddr + 1];
11      ReadDataBlock[95:64] <= DRAM[DataAddr + 2];
12      ReadDataBlock[127:96] <= DRAM[DataAddr + 3];
13    end
14    Ready <= 1;
15  end else begin
16    Ready <= 0;
17  end

```

18 | end

Listing 8: DRAM Block Read/Write Logic

2.2.5 Module: Arbiter

Behavior: The **Arbiter** module grants bus ownership among the 4 L1 caches using round-robin priority. It preserves ownership while the current owner continues requesting and selects the next requester when the bus is released.

Key signals:

- **BusReq[3:0]**: bus requests from L1-0..L1-3
- **BusGrant[3:0]**: one-hot bus grant output
- **BusBusy**: shared busy line (tri-stated by arbiter, driven by bus users)

What is Round Robin? It's an arbitration scheme that cycles through requesters in a fixed order, granting access to the next requester in line after the current one finishes. This scheme is also used in various other places regarding computers, such as process scheduling in operating systems.

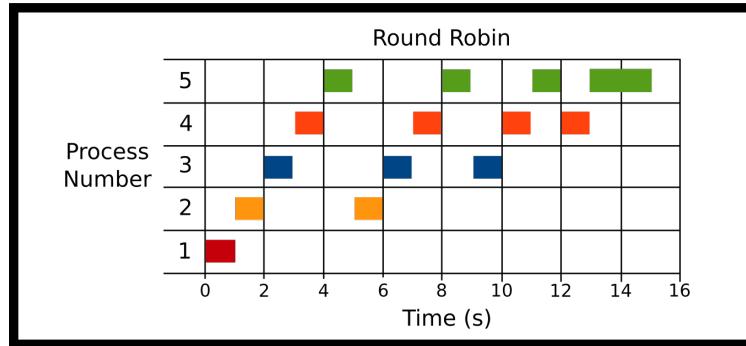


Figure 3: Round-robin arbitration illustration (adapted from [3]).

Essentially it loops through the requesters in a circle.

Key State:

- **grant_pointer**: round-robin starting priority pointer
- **grant_reg/grant_next**: current and next one-hot grant state

Implementation:

```

1 always_comb begin
2   grant_next = grant_reg;
3
4   if (!grant_reg) begin
5     if ((grant_reg & BusReq) == 4'b0000) begin
6       grant_next = 4'b0000;
7     end
8   end
9 
```

```

10   if (grant_next == 4'b0000) begin
11     for (int i = 0; i < 4; i = i + 1) begin
12       if (BusReq[(grant_pointer + i[1:0]) & 2'b11] && grant_next == 4'
13         b0000)
14         grant_next[(grant_pointer + i[1:0]) & 2'b11] = 1'b1;
15     end
16   end

```

Listing 9: Round-Robin Grant Selection

```

1 always_ff @(posedge clk) begin
2   if (reset) begin
3     grant_pointer <= 2'b00;
4     grant_reg <= 4'b0000;
5   end else begin
6     grant_reg <= grant_next;
7     if (grant_reg == 4'b0000 && grant_next != 4'b0000) begin
8       for (int i = 0; i < 4; i = i + 1)
9         if (grant_next[i]) grant_pointer <= i[1:0] + 2'b01;
10    end
11  end
12 end

```

Listing 10: Grant Register and Pointer Update

3 Verification and Results

3.0.1 Purpose

There are 5 main tests that are being accounted for in this implementation:

- Single processor reads
- Single processor writes
- Cache coherence
- Arbitration
- Memory eviction

These tests inadvertently cover a lot of different scenarios, such as:

- Cache hits and misses
- Writebacks of dirty blocks
- Coherence state transitions (e.g. SHARED to MODIFIED on write)
- Bus transactions (BusRd, BusRdX, BusUpgr)
- DRAM interactions on L2 misses

This state transition diagram is what our MESI model should be following:

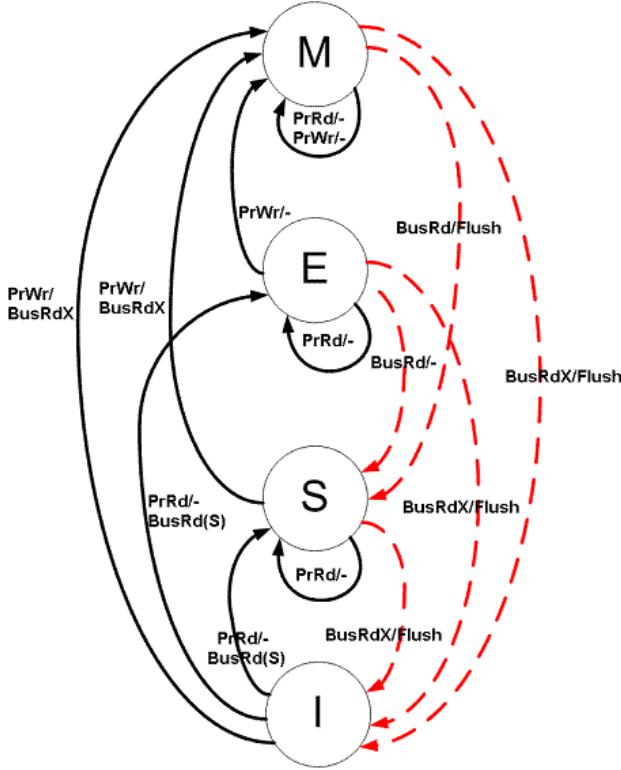


Figure 4: MESI state transition diagram (from [1], credited to [2]).

3.1 Test 1: Memory Read

This test ensures that memory is being loaded correctly into each cache block. *Stimulus file: Listing 11 (Appendix A).*

Table 1: Test 1 (Memory Read): stimulus and expected cache/coherence behavior

Cycle	Address	Op	Req.	Expected behavior
0	0x00000000	Read	L1-0	Cold miss; BusRd; fill line {0x00,0x04,0x08,0xC}.
1	0x00000004	Read	L1-0	Same line as cycle 0; hit (offset change only).
2	0x00000008	Read	L1-0	Same line; hit.
3	0x0000000C	Read	L1-0	Same line; hit.
4	0x00000100	Read	L1-1	New line for L1-1; miss and fill expected.
5	0x00000104	Read	L1-2	Same line family as cycle 4; shared copy behavior expected.
6	0x00000108	Read	L1-3	Same shared line family; shared-state behavior remains valid.

3.1.1 Waveforms

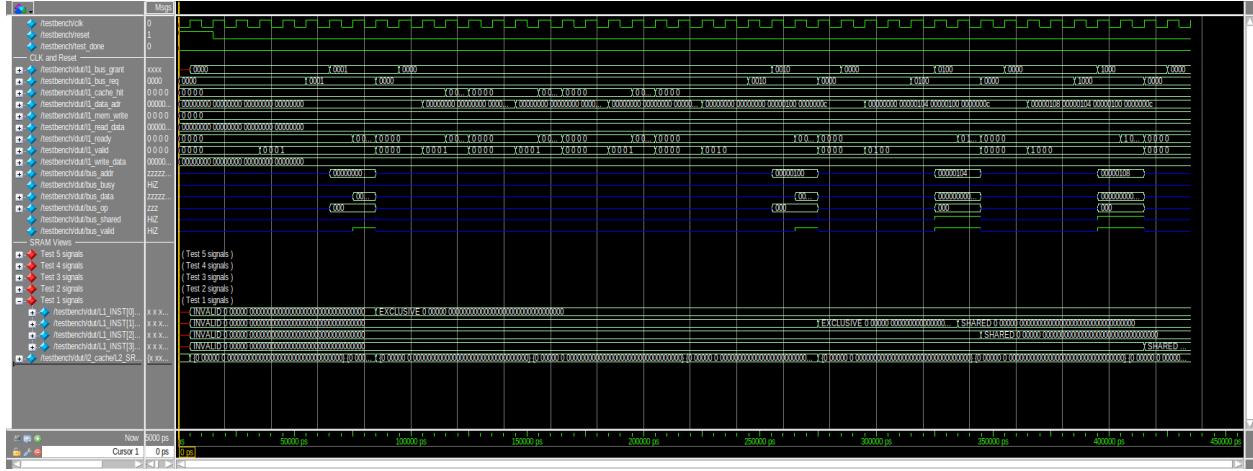


Figure 5: Test 1 waveform showing memory read operations, cache hits/misses, and MESI state transitions.

Given the waveforms, there are some important things to note to know that the test is working correctly:

- `l1_cache_hit` is 0 for the first request, but changes values for the next 3 requests. It's not visible in the image, but it does change its value to 1 for the next 3 requests, which means that those are hits.
- The MESI state for the first request changes to EXCLUSIVE, and then changes to SHARED for the next 3 requests, which is expected since the first request is a cold miss, and the next 3 requests are hits to the same line, which should be in SHARED state since multiple caches have a copy of the same line.
- The bus transactions for the first request is a BusRd, which is expected since it's a miss. The next 3 requests don't have any bus transactions, which is expected since they are hits.

3.2 Test 2: Memory Write

This test checks single-core and multi-core write behavior, including write misses, ownership acquisition, and dirty-line updates. *Stimulus file: Listing 12 (Appendix A)*.

Table 2: Test 2 (Memory Write): stimulus and expected cache/coherence behavior

Cycle	Address	Op	Req.	Data	Expected behavior
0	0x00000010	Write	L1-0	0xDEADBEEF	Write miss; BusRdX (or write-allocate path); line filled, word updated, line becomes MODIFIED.
1	0x00000014	Write	L1-0	0xBEDFACED	Same cache line family as prior access; write hit in owned line, remains MODIFIED.
2	0x00000020	Write	L1-1	0xAAAABBBB	Independent write by another core; miss + ownership acquisition expected for L1-1 line.
3	0x00000024	Write	L1-2	0xCCCCDDDD	Miss for L1-2 on its line; BusRdX/fill/update expected.
4	0x00000028	Write	L1-3	0xEEEEAAAA	Miss for L1-3 on its line; data written and line ends MODIFIED.

3.2.1 Waveforms



Figure 6: Test 2 waveform showing memory write operations, cache hits/misses, MESI state transitions, and bus transactions.

- First write to each new line should show miss behavior (bus transaction present).
- Re-write within an already-owned line should avoid refill and complete as a hit.
- Updated word lane in cache line should match written data values.

3.3 Test 3: Cache Coherence (Snooping + MESI Transitions)

This test validates snoop-driven coherence behavior across L1 caches, including SHARED formation, invalidation on write, and read-after-invalidate correctness. *Stimulus file: Listing 13 (Appendix A).*

Table 3: Test 3 (Cache Coherence): stimulus and expected cache/coherence behavior

Cycle	Address	Op	Req.	Data	Expected behavior
0	0x00000200	Read	L1-0	0x00000000	Concurrent read with L1-1 to same line; first reader may fetch, both should converge to SHARED.
0	0x00000200	Read	L1-1	0x00000000	Snoop/BusShared path active; both caches hold coherent shared copies after fill/response.
1	0x00000204	Read	L1-0	0x00000000	L1-0 reads line first (E/S depending on sharers).
2	0x00000204	Write	L1-1	0xFFFFFFFF	L1-1 write requires ownership (Bus-RdX/BusUpgr); peer shared copy invalidated; L1-1 becomes MODIFIED.
3	0x00000204	Read	L1-0	0x00000000	L1-0 re-read after invalidation must miss/refetch; returned data should reflect L1-1 write.
4	0x00000208	Read	L1-0	0x00000000	Concurrent read pair with L1-1 creates shared line in both caches.
4	0x00000208	Read	L1-1	0x00000000	Shared-state confirmation via snoop-visible bus behavior.
5	0x00000208	Write	L1-0	0x99999999	Writer upgrades/gets ownership; other sharer invalidated; writer ends MODIFIED.

3.3.1 Waveforms

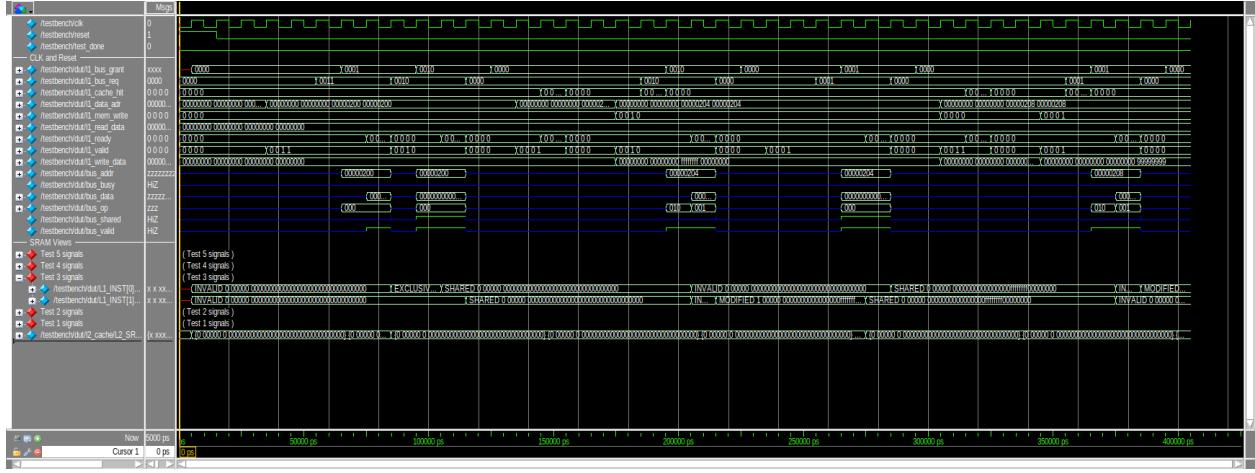


Figure 7: Test 3 waveform showing coherence snoops, shared-line formation, invalidation, and ownership transfer.

- Concurrent reads to the same address should show shared-copy behavior (BusShared/snoop-visible response).
- Write from one cache to a shared line should invalidate peer copies and transition writer to MODIFIED.
- Read-after-invalidate should miss and return updated data.

3.4 Test 4: Arbitration (Round-Robin Fairness Under Contention)

This test verifies that simultaneous requests from all four L1 caches are granted in fair round-robin order for both read and write bursts. *Stimulus file: Listing 14 (Appendix A)*.

Table 4: Test 4 (Arbitration): concurrent request groups and expected arbiter behavior

Cycle	Address	Op	Req.	Data	Expected behavior
0	0x000000300	Read	L1-0	0x00000000	All four reads arrive together; grants should serialize in round-robin order without starvation.
0	0x000000304	Read	L1-1	0x00000000	Request remains pending until granted; Bus-Grant should be one-hot each service window.
0	0x000000308	Read	L1-2	0x00000000	Arbiter rotates priority pointer after each completed ownership window.
0	0x00000030C	Read	L1-3	0x00000000	Final requester in first contention group must still be serviced correctly.
10	0x000000310	Write	L1-0	0xAFFFFFFF	Second contention group (writes): same fairness behavior expected under write traffic.
10	0x000000314	Write	L1-1	0xBFFFFFFF	One-hot grant + no overlap on bus ownership.
10	0x000000318	Write	L1-2	0xCFFFFFFF	Each requester eventually receives grant; no live-lock/starvation.
10	0x00000031C	Write	L1-3	0xDFFFFFFF	Rotation should continue consistently from prior pointer state.

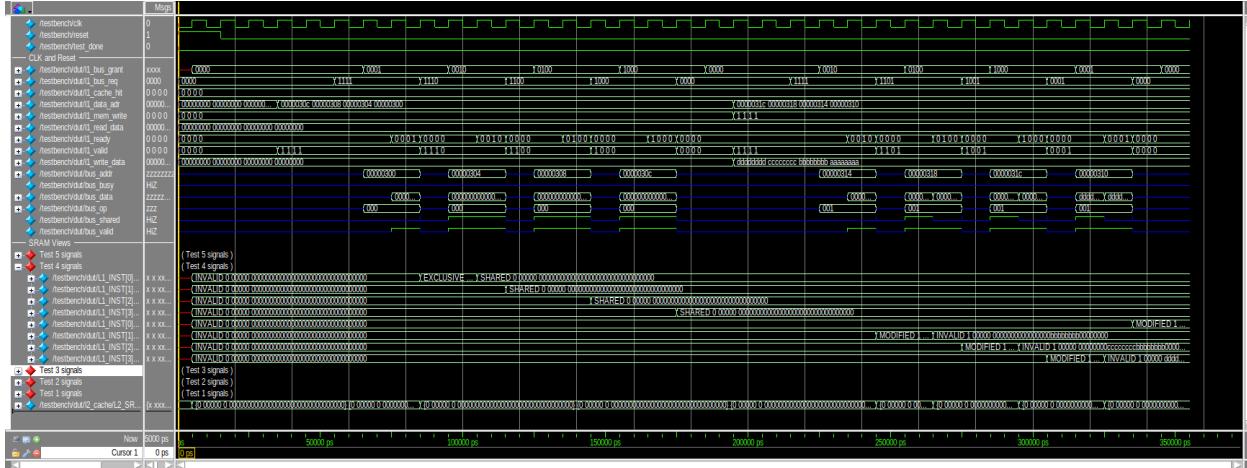


Figure 8: Test 4 waveform showing concurrent requests and round-robin bus grant sequencing.

3.4.1 Waveforms

- For each contention burst, **BusGrant [3:0]** should be one-hot and rotate fairly.
- **BusBusy** should remain asserted while an owner is active, then release for next grant.
- All four requesters in each burst should complete with no starvation.

3.5 Test 5: Memory Eviction and Writeback

This test validates direct-mapped conflict eviction, dirty-line writeback, and correct data recovery after eviction. *Stimulus file: Listing 15 (Appendix A).*

Table 5: Test 5 (Eviction/Writeback): stimulus and expected behavior

Cycle	Address	Op	Req.	Data	Expected behavior
0	0x00000000	Write	L1-0	0xAFFFFFFF	Write miss/fill then update; line in L1-0 becomes MODIFIED (dirty).
1	0x00002000	Write	L1-0	0xBFFFFFFF	Same index, different tag: dirty victim eviction; writeback expected before/with replacement.
2	0x00004000	Write	L1-0	0xCFFFFFFF	Another conflict on same index: second eviction/writeback path exercised.
3	0x00000000	Read	L1-0	0x00000000	Re-read original line should miss (evicted) and fetch most recent committed value from lower level.
4	0x00006000	Write	L1-1	0xDFFFFFFF	Repeat eviction behavior on a different core; verifies per-core buffer/control path.
5	0x00008000	Write	L1-1	0xEFFFFFFF	Conflict replacement in L1-1 triggers dirty handling consistently.

3.5.1 Waveforms

- Conflict addresses mapping to the same index should trigger eviction/replacement.
- Dirty victim lines should generate writeback traffic before data loss.
- Post-eviction read should miss and return the correct committed value.

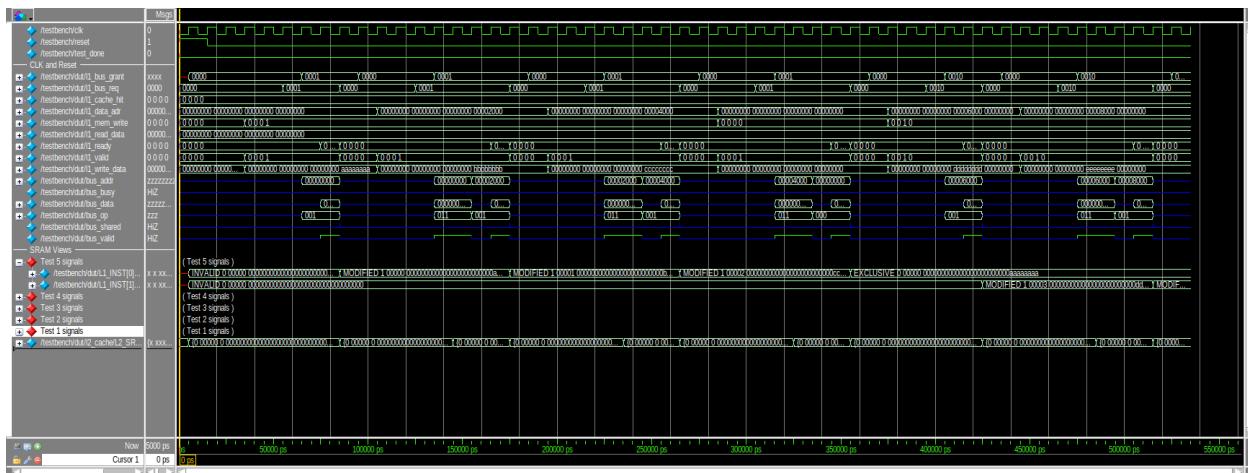


Figure 9: Test 5 waveform showing conflict misses, dirty evictions, writeback transactions, and reload behavior.

- It’s unable to see the writeback transactions in the waveform, but they do happen in the simulation, which can be verified by looking at the L2 and DRAM interactions, and how the L2 line changes at the evicting address.

A Test Vector Files

A.1 Test 1 Vectors (Memory Read)

```
1 // Test 1: Single reads from different addresses
2 // execution each read completes before next starts
3 // ADDRESS_MEMWRITE_L1ID_DATA_CYCLE
4 00000000_0_0_00000000_0
5 00000004_0_0_00000000_1
6 00000008_0_0_00000000_2
7 0000000C_0_0_00000000_3
8 00000100_0_1_00000000_4
9 00000104_0_2_00000000_5
10 00000108_0_3_00000000_6
```

Listing 11: Test 1 vectors: test1.single_read.txt

A.2 Test 2 Vectors (Memory Write)

```
1 // Test 2: Single writes with 32-bit data
2 // Each write completes before next one starts (different cycles)
3 // ADDRESS_MEMWRITE_L1ID_DATA_CYCLE
4 00000010_1_0_DEADBEEF_0
5 00000014_1_0_BEDFACED_1
6 00000020_1_1_AAAABBBA_2
7 00000024_1_2_CCCCCDDD_3
8 00000028_1_3_EEEEAAAA_4
```

Listing 12: Test 2 vectors: test2.single_write.txt

A.3 Test 3 Vectors (Cache Coherence)

```
1 // Test 3: Coherence scenario - snooping initiated MESI state transitions
2 // ADDRESS_MEMWRITE_L1ID_DATA_CYCLE
3 // Scenario 1 (Cycle 0): Both L1s read same address (both should reach
4 // SHARED state)
5 // Concurrent reads on same cycle to test snoop response
6 00000200_0_0_00000000_0
7 00000200_0_1_00000000_0
8 // Scenario 2 (Cycle 1-3): L1-0 reads, L1-1 writes (upgrading invalidates
9 // others)
10 // Sequential to show state transition
11 00000204_0_0_00000000_1
12 00000204_1_1_FFFFFFFF_2
13 00000204_0_0_00000000_3
14 // Scenario 3 (Cycle 4-5): Multi-L1 coherence - read-then-write pattern
15 // Both read concurrently, then L1-0 writes (invalidates L1-1's copy)
16 00000208_0_0_00000000_4
17 00000208_0_1_00000000_4
18 00000208_1_0_99999999_5
```

Listing 13: Test 3 vectors: test3.coherence.txt

A.4 Test 4 Vectors (Arbitration)

```

1 // Test 4: Arbitration - concurrent requests from all 4 L1s
2 // Format: ADDRESS_MEMWRITE_L1ID_DATA_CYCLE
3 //
4 // KEY TIMING CONCEPT:
5 // - Cycle 0: All 4 reads issued SIMULTANEOUSLY (concurrent)
6 //   The arbiter grants them one-by-one in round-robin order
7 // - Cycle 10: All 4 writes issued SIMULTANEOUSLY (concurrent, delayed)
8 //
9 // Round 1 (Cycle 0): CONCURRENT reads - arbiter tests round-robin
10 00000300_0_0_00000000_0
11 00000304_0_1_00000000_0
12 00000308_0_2_00000000_0
13 0000030C_0_3_00000000_0
14 // Round 2 (Cycle 10): CONCURRENT writes - arbiter tests round-robin
15 00000310_1_0_AAAAAAAA_10
16 00000314_1_1_BBBBBBBB_10
17 00000318_1_2_CCCCCCCC_10
18 0000031C_1_3_DDDDDDDD_10

```

Listing 14: Test 4 vectors: test4.arbitration.txt

A.5 Test 5 Vectors (Eviction/Writeback)

```

1 // Test 5: Cache line eviction and write-back
2 // ADDRESS_MEMWRITE_L1ID_DATA_CYCLE
3 // Addresses with same index but different tags (direct-mapped cache)
4 // This tests write-back buffer and dirty line eviction
5 // SEQUENTIAL execution: each write must complete before next starts
6
7 // Step 1 (Cycle 0): Write to address 0x00000000
8 00000000_1_0_AAAAAAAA_0
9
10 // Step 2 (Cycle 1): Write to address 0x00002000 (different tag, same
11 //   index, evicts prev)
11 00002000_1_0_BBBBBBBB_1
12
13 // Step 3 (Cycle 2): Write to address 0x00004000 (evicts second line)
14 00004000_1_0_CCCCCCCC_2
15
16 // Step 4 (Cycle 3): Read from address 0x00000000 (should miss, line was
17 //   evicted)
17 00000000_0_0_00000000_3
18
19 // Step 5 (Cycle 4-5): Test eviction with different L1s
20 00006000_1_1_DDDDDDDD_4
21 00008000_1_1_EEEEEEEE_5

```

Listing 15: Test 5 vectors: test5.eviction.txt

B Reproducibility Notes

This report and all results were generated using the following toolchain:

- Quartus Prime 20.1.1
- ModelSim - Intel FPGA Starter Edition 2020.1
- Linux environment (project root: `project1/`)

To reproduce the verification results:

- Open `cache_coherence.mpf` in ModelSim
- Ensure that all the files there are `.v` files attached their corresponding module
- Start a simulation with `testbench` as the top-level module
- When in the waveform viewer, go to `File > Load > Macro File...` and select `simulation/wave.do` to load the relevant signals for viewing
- `Simulate > Run > Run-All`

References

- [1] Wikipedia contributors, *MESI protocol*, https://en.wikipedia.org/wiki/MESI_protocol, accessed Feb. 2026.
- [2] D. E. Culler and J. P. Singh, *Parallel Computer Architecture: A Hardware/Software Approach*, Morgan Kaufmann, 1997.
- [3] University of Illinois CS 341 course material, *Round Robin Scheduling*, <https://cs341.cs.illinois.edu/>, accessed Feb. 2026.
- [4] University of Waterloo (CS 450), *MESI Handout*, <https://student.cs.uwaterloo.ca/~cs450/w18/public/mesiHandout.pdf>, accessed Feb. 2026.
- [5] D. J. Sorin, M. D. Hill, and D. A. Wood, *A Primer on Memory Consistency and Cache Coherence*, Morgan & Claypool Publishers, 2011.